

# Parallelizing NN for sound classification

HPPL Project  
Svetlana Pavlova

# Motivation

**Neural Networks (NNs)** have proven very effective in image classification tasks, which gave rise to the design of various architectures.

NNs achieve state of the art results on image classification tasks and offer a variety of ready to use pre trained backbones.

Instead of directly using the sound file as an amplitude vs time signal it is wished to convert the audio signal into an image - spectrogram.

# Data overview – ESC-50 dataset

Animals	Natural soundscapes & water sounds	Human, non-speech sounds	Interior/domestic sounds	Exterior/urban noises
Dog	Rain	Crying baby	Door knock	Helicopter
Rooster	Sea waves	Sneezing	Mouse click	Chainsaw
Pig	Crackling fire	Clapping	Keyboard typing	Siren
Cow	Crickets	Breathing	Door, wood creaks	Car horn
Frog	Chirping birds	Coughing	Can opening	Engine
Cat	Water drops	Footsteps	Washing machine	Train
Hen	Wind	Laughing	Vacuum cleaner	Church bells
Insects (flying)	Pouring water	Brushing teeth	Clock alarm	Airplane
Sheep	Toilet flush	Snoring	Clock tick	Fireworks
Crow	Thunderstorm	Drinking, sipping	Glass breaking	Hand saw

2000 environmental audio recordings

5-second-long recordings


50 semantical classes (with 40 examples per class)

# Data overview – waveplots

cat



crackling\_fire




clapping



dog



thunderstorm



footsteps




# Mel scale

Humans do not perceive frequencies on a linear scale. We are better at detecting differences in lower frequencies than higher frequencies.


For example, we can easily tell the difference between 500 and 1000 Hz, but we will hardly be able to tell a difference between 10,000 and 10,500 Hz, even though the distance between the two pairs are the same.

In 1937, Stevens, Volkman, and Newmann proposed a unit of pitch such that equal distances in pitch sounded equally distant to the listener. This is called the **mel scale**.




# From wav to melspectrograms

Using librosa library to convert wav to melspectrogram and then convert melspectrogram to image (mapped the y-axis (frequency) onto the mel scale to form the melspectrogram)



clock tick



# From wav to melspectrograms – useful functions

```
def get_melspectrogram_db(file_path, sr=None, n_fft=2048, hop_length=512, n_mels=128, fmin=20, fmax=8300, top_db=80):  
    wav, sr = librosa.load(file_path, sr=sr)  
    if wav.shape[0] < 5*sr:  
        wav = np.pad(wav, int(np.ceil((5*sr - wav.shape[0])/2)), mode='reflect')  
    else:  
        wav = wav[:5*sr]  
    spec = librosa.feature.melspectrogram(wav, sr=sr, n_fft=n_fft,  
                                         hop_length=hop_length, n_mels=n_mels, fmin=fmin, fmax=fmax)  
    spec_db = librosa.power_to_db(spec, top_db=top_db)  
    return spec_db
```

# From wav to melspectrograms – useful functions

```
def spec_to_image(spec, eps=1e-6):  
    mean = spec.mean()  
    std = spec.std()  
    spec_norm = (spec - mean) / (std + eps)  
    spec_min, spec_max = spec_norm.min(), spec_norm.max()  
    spec_scaled = 255 * (spec_norm - spec_min) / (spec_max - spec_min)  
    spec_scaled = spec_scaled.astype(np.uint8)  
    return spec_scaled
```



# Next steps for NN building

1. **Loading data in Pytorch** - build dataloaders to preprocess and load data
2. **Building Model** - use custom model
3. **Training** - CrossEntropyLoss and Adam optimizer are used. The model is trained for 20 epochs with stable learning rate

# Train a neural network on a GPU

**CUDA** (Compute Unified Device Architecture) is a parallel computing platform developed by NVIDIA for general computing on graphics processing units (GPUs). With CUDA, developers can accelerate computing applications by leveraging the power of GPUs.

Let's first define our GPU as the first visible cuda device.

```
if torch.cuda.is_available():  
    device=torch.device('cuda')  
else:  
    device=torch.device('cpu')
```

# Train a neural network on a GPU

Send the network to the GPU:


```
1 model = ESC50Model(input_shape=(1,128,431), batch_size=16, num_cats=50).to(device)
```

Send data at each step to the GPU:

```
x = x.to(device, dtype=torch.float32)  
y = y.to(device, dtype=torch.long)
```

# Runtime comparison

CPU:



With the use of GPU the network training lasted about 12 minutes. The same training on a regular processor will last for 3 hours. The difference is significant, this is because our network is big. When using large arrays for training, the difference between the speed of a GPU and a traditional processor increases.

GPU:




# Final model results

Valid-Accuracy : 0.545

dog.wav [ <=> ] 3.61M 4.89MB/s in 0.7s

2021-12-20 03:48:42 (4.89 MB/s) - 'dog.wav' saved [3786304]


dog



cat.wav [ <=> ] 873.05K 1.92MB/s in 0.4s

2021-12-20 03:33:42 (1.92 MB/s) - 'cat.wav' saved [894000]

laughing



# Project results

- Learned how to use NN's
- Learned about sound processing
- Learned about image classification with NN's
- Learned how to use GPU to speed up model training

## Improvements

- Writing own parallelization scripts instead of using Pytorch distributed package
- Think about better NN's model architecture